

N TLA+ Specification

MODULE *Nezha*

EXTENDS *Naturals*, *TLC*, *FiniteSets*, *Sequences*

Bounds for Model Check [Configurable]

Time Range [Configurable]
 $MaxTime \triangleq 3$

Each client is only allowed to submit $MaxReqNum$ requests [Configurable]
In the specification, we will only consider two roles, client and replicas
(i.e. it can be considered as co-locating one proxy with one client)
For the proxy-based design, we just need to replace client with proxy,
and then the specification describes the interaction between proxy and replicas
 $MaxReqNum \triangleq 1$

The leader is only allowed to crash when the view $< MaxViews$ [Configurable]
 $MaxViews \triangleq 3$

These variables are used to implement at-most-once primitives
i.e. The variables record the messages processed by *Replicas/Clients*, so
that the *Replicas/Clients* will not process twice

VARIABLE $vReplicaProcessed$, Messages that have been processed by replicas
 $vClientProcessed$ Messages that have been processed by clients

VARIABLE *DebugAction*

Constants

The set of replicas and an ordering of them

CONSTANTS *Replicas*, *ReplicaOrder*, *Clients*, *LatencyBounds*
ASSUME $IsFiniteSet(Replicas)$
ASSUME $ReplicaOrder \in Seq(Replicas)$

$F \triangleq (Cardinality(Replicas) - 1) \div 2$
 $ceilHalfF \triangleq IF (F \div 2) * 2 = F THEN F \div 2 ELSE (F + 1) \div 2$
 $floorHalfF \triangleq F \div 2$
 $QuorumSize \triangleq F + 1$
 $FastQuorumSize \triangleq F + ceilHalfF + 1$
 $RecoveryQuorumSize \triangleq ceilHalfF + 1$
 $FastQuorums \triangleq \{R \in SUBSET (Replicas) : Cardinality(R) \geq FastQuorumSize\}$
 $Quorums \triangleq \{R \in SUBSET (Replicas) : Cardinality(R) * 2 > Cardinality(Replicas)\}$

Replica Statuses
 $StNormal \triangleq 1$

$StViewChange \triangleq 2$
 $StRecovering \triangleq 3$

Message Types

$MClientRequest \triangleq 1$ Sent by client to replicas
 $MFastReply \triangleq 2$ Fast Reply Message
 $MSlowReply \triangleq 3$ Slow Reply Message
 $MLogIndex \triangleq 4$ *LogIndex*
 $MLogEntry \triangleq 5$ Log entry, different from index, it includes command field, which can be large in practice
 $MIndexSync \triangleq 6$ Sync message during the index sync process
 $MMissEntryRequest \triangleq 7$ Sent by followers once they fail to find the entry on itself
 $MMissEntryReply \triangleq 8$ Response to *MMissEntryRequest*, providing the missing entries

 $MViewChangeReq \triangleq 9$ Sent when leader/sequencer failure detected
 $MViewChange \triangleq 10$ Sent to *ACK* view change
 $MStartView \triangleq 11$ Sent by new leader to start view

The following messages are mainly used for periodic sync

Just as described in *NOPaxos*, it is an optional optimization to enable fast recovery after failure

$MSyncPrepare \triangleq 12$ Sent by the leader to ensure *log* durability
 $MSyncRep \triangleq 13$ Sent by followers as *ACK*
 $MSyncCommit \triangleq 14$ Sent by leaders to indicate stable *log*

The following messages are mainly used for replica recovery

$MCrashVectorReq \triangleq 15$
 $MCrashVectorRep \triangleq 16$
 $MRecoveryReq \triangleq 17$
 $MRecoveryRep \triangleq 18$
 $MStateTransferReq \triangleq 19$
 $MStateTransferRep \triangleq 20$

Message Schemas

$ViewIDs \triangleq [leaderNum \mapsto n \in (1 \dots)]$

$\backslash *$ $\langle clientID, requestID \rangle$ uniquely identifies one request on one replica
 $\backslash *$ But across replicas, the same $\langle clientID, requestID \rangle$ may have different deadlines
 $\backslash *$ (the leader may modify the deadline to make the request eligible to enter the early-buffer)
 $\backslash *$ so $\langle deadline, clientID, reqID \rangle$ uniquely identifies one request across replicas

ClientRequest

$[$ $mtype \mapsto MClientRequest,$
 $sender \mapsto c \in Clients,$
 $dest \mapsto r \in Replicas,$
 $requestID \mapsto i \in (1 \dots),$
 $command \mapsto "",$
 $s \mapsto t \in (1 \dots MaxTime),$
 $l \mapsto l \in (1 \dots MaxBound)$
 $]$

```

\ * logSlotNum is not necessary and it is not described in the paper
\ * Here we include logSlotNum in FastReply and SlowReply messages
\ * to facilitate the check of Linearizability invariant

```

FastReply

```

[ mtype      ↦ MFastReply,
  sender     ↦ r ∈ Replicas,
  dest       ↦ c ∈ Clients,
  viewID     ↦ v ∈ ViewIDs,
  requestID  ↦ i ∈ (1 .. vClientReqNum)
  hash       ↦ [ log ↦ vLogs[1 .. n],
                  cv ↦ crashVector
                ]
  deadline   ↦ i ∈ (1 .. MaxTime + MaxBound),
  logSlotNum ↦ n ∈ (1 .. )
]

```

SlowReply

```

[ mtype      ↦ MSlowReply,
  sender     ↦ r ∈ Replicas,
  dest       ↦ c ∈ Clients,
  viewID     ↦ v ∈ ViewIDs,
  requestID  ↦ i ∈ (1 .. vClientReqNum)
  logSlotNum ↦ n ∈ (1 .. )
]

```

LogIndex

```

[ mtype      ↦ MLogIndex,
  clientID   ↦ c ∈ Clients,
  requestID  ↦ i ∈ (1 .. vClientReqNum),
  deadline   ↦ i ∈ (1 .. MaxTime + MaxBound),
]

```

LogEntry

```

[ mtype      ↦ MLogEntry,
  clientID   ↦ c ∈ Clients,
  requestID  ↦ i ∈ (1 .. vClientReqNum),
  deadline   ↦ i ∈ (1 .. MaxTime + MaxBound),
  command    ↦ ""
]

```

IndexSync

```

[ mtype      ↦ MIndexSync,
  sender     ↦ r ∈ Replicas,
  dest       ↦ c ∈ Clients,
  viewID     ↦ v ∈ ViewIDs,
  logIndices ↦ index ∈ vLogs[leaderIdx]
]

```

MMissEntryRequest

```

[ mtype      ↦ MMissEntryRequest,
  sender     ↦ r ∈ Replicas,
  dest       ↦ d ∈ Replicas,
]

```

```

    viewID    ↦ v ∈ ViewIDs,
    miss      ↦ {log indices}
  ]

MMissEntryRequest
[ mtype      ↦ MMissEntryReply,
  sender     ↦ r ∈ Replicas,
  dest       ↦ d ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  entries    ↦ {log entries}
]

ViewChangeReq
[ mtype ↦ MViewChangeReq,
  sender ↦ r ∈ Replicas,
  dest  ↦ r ∈ Replicas,
  viewID ↦ v ∈ ViewIDs,
  cv    ↦ crash vector
]

ViewChange
[ mtype      ↦ MViewChange,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  lastNormal ↦ v ∈ ViewIDs,
  log        ↦ l ∈ vLogs[1 .. n],
  cv         ↦ crash vector
]

StartView
[ mtype      ↦ MStartView,
  dest       ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  log        ↦ l ∈ vLogs[1 .. n],
  cv         ↦ crash vector
]

SyncPrepare
[ mtype      ↦ MSyncPrepare,
  dest       ↦ r ∈ Replicas,
  sender     ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  log        ↦ l ∈ vLogs[1 .. n] ]

SyncRep
[ mtype      ↦ MSyncRep,
  dest       ↦ r ∈ Replicas,
  sender     ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  logSlotNumber ↦ n ∈ (1 .. ) ]

SyncCommit

```

```

[ mtype      ↦ MSyncCommit,
  dest       ↦ r ∈ Replicas,
  sender     ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  log        ↦ l ∈ vLogs[1 .. n] ]

CrashVectorReq
[ mtype      ↦ MCrashVectorReq,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  nonce      ↦ nonce
]

CrashVectorRep
[ mtype      ↦ MCrashVectorRep,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  nonce      ↦ nonce,
  cv         ↦ vector of counters
]

RecoveryReq
[ mtype      ↦ MRecoveryReq,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  cv         ↦ vector of counters
]

RecoveryRep
[ mtype      ↦ MRecoveryRep,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  cv         ↦ vector of counters
]

StateTransferReq
[ mtype      ↦ MStateTransferReq,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  cv         ↦ vector of counters
]

StateTransferRep
[ mtype      ↦ MStateTransferRep,
  sender     ↦ r ∈ Replicas,
  dest       ↦ r ∈ Replicas,
  viewID     ↦ v ∈ ViewIDs,
  log        ↦ l ∈ vLogs[1 .. n] ],
  cv         ↦ vector of counters
]

```

Variables

Network State

VARIABLE *messages* Set of all messages sent

$networkVars \triangleq \langle messages \rangle$
 $InitNetworkState \triangleq messages = \{\}$

Used as a dummy value

$NULLLog \triangleq [\begin{array}{ll} deadline & \mapsto 0, \\ clientID & \mapsto 0, \\ requestID & \mapsto 0 \end{array}]$

Replica State

VARIABLES *vLog*,

vEarlyBuffer,

vReplicaStatus,

vViewID,

vReplicaClock,

vLastNormView,

vViewChanges,

vSyncPoint,

vLateBuffer,

vTentativeSync,

vSyncReps,

vCommitPoint,

vUUIDCounter,

vCrashVector,

vCrashVectorReps,

vRecoveryReps

Log of values

The early buffer to hold request,
and release it after clock passes its deadline ($s + l$)

One of *StNormal*, *StViewChange*, *StRecovering*

Current *viewID* replicas recognize

Current Time of the replica

Last views in which replicas had status *StNormal*

Used for logging view change votes

Latest synchronization point,
to which the replica state (*vLog*) is consistent with the leader.

The late buffer Used to store the requests
which are not eligible to enter *vEarlyBuffer*

Used by leader to mark current *syncPrepare* point (during periodic sync process)
(Actually, *vSyncPoint* and *vTentativeSync* can be merged into one *Var*
However, we decouple them to make the spec easy to understand)

Used for logging sync reps at leader

Different from *vSyncPoint*,
vCommitPoint indicates that the logs before this point has been replicated to majority
So followers can safely execute requests (*log* entries) up to *vCommitPoint*
Refer to “Acceleration of Recovery“ para in *Sec 6*

Locally unique string (for *CrashVectorReq*)

CrashVector, initialized as all-zero vector

CrashVectorRep Set

RecoveryRep Set

$replicaVars \triangleq \langle vLog, vEarlyBuffer, \\ vViewID, vReplicaClock, \\ vLastNormView, vViewChanges, vReplicaStatus, \\ vSyncPoint, vLateBuffer, \\ vTentativeSync, vSyncReps, vCommitPoint, \\ vUUIDCounter, vCrashVector, \\$

$vCrashVectorReps, vRecoveryReps\rangle$

$InitReplicaState \triangleq$

- $\wedge vLog = [r \in Replicas \mapsto \langle \rangle]$
- $\wedge vEarlyBuffer = [r \in Replicas \mapsto \{\}]$
- $\wedge vViewID = [r \in Replicas \mapsto 1]$ 0 should also be okay
- $\wedge vReplicaClock = [r \in Replicas \mapsto 1]$
- $\wedge vLastNormView = [r \in Replicas \mapsto 1]$
- $\wedge vViewChanges = [r \in Replicas \mapsto \{\}]$
- $\wedge vReplicaStatus = [r \in Replicas \mapsto StNormal]$
- $\wedge vSyncPoint = [r \in Replicas \mapsto 0]$
- $\wedge vLateBuffer = [r \in Replicas \mapsto \{\}]$
- $\wedge vTentativeSync = [r \in Replicas \mapsto 0]$
- $\wedge vSyncReps = [r \in Replicas \mapsto \{\}]$
- $\wedge vCommitPoint = [r \in Replicas \mapsto 0]$
- $\wedge vCrashVector = [r \in Replicas \mapsto [rr \in Replicas \mapsto 0]]$
- $\wedge vCrashVectorReps = [r \in Replicas \mapsto \{\}]$
- $\wedge vRecoveryReps = [r \in Replicas \mapsto \{\}]$
- $\wedge vUUIDCounter = [c \in Replicas \mapsto 0]$

Client State

VARIABLES $vClientClock,$ Current Clock Time of the client
 $vClientReqNum$ The number of requests that have been sent by this client

$InitClientState \triangleq$

- $\wedge vClientClock = [c \in Clients \mapsto 1]$
- $\wedge vClientReqNum = [c \in Clients \mapsto 0]$

$clientVars \triangleq \langle vClientClock, vClientReqNum \rangle$

Set of all vars

$vars \triangleq \langle networkVars, replicaVars, clientVars \rangle$

\ * Initial state

$Init \triangleq$

- $\wedge InitNetworkState$
- $\wedge InitReplicaState$
- $\wedge InitClientState$
- $\wedge vReplicaProcessed = [r \in Replicas \mapsto \{\}]$
- $\wedge vClientProcessed = [c \in Clients \mapsto \{\}]$
- $\wedge DebugAction = \langle "Init", "" \rangle$

Helpers

$NumofReplicas(status) \triangleq Cardinality(\{r \in Replicas : vReplicaStatus[r] = status\})$

$DuplicateRep(ReplySet, m) \triangleq m.sender \in \{mm.sender : mm \in ReplySet\}$

$Pick(S) \triangleq \text{CHOOSE } s \in S : \text{TRUE}$

Convert a Set to Sequence

RECURSIVE $Set2Seq(-)$
 $Set2Seq(S) \triangleq \text{IF } Cardinality(S) = 0 \text{ THEN } \langle \rangle$
 ELSE
 LET
 $x \triangleq \text{CHOOSE } x \in S : \text{TRUE}$
 IN
 $\langle x \rangle \circ Set2Seq(S \setminus \{x\})$

Convert a Sequence to Set

$Seq2Set(seq) \triangleq \{seq[i] : i \in \text{DOMAIN } seq\}$

$Max(S) \triangleq \text{CHOOSE } x \in S : \forall y \in S : x \geq y$

$Min(S) \triangleq \text{CHOOSE } x \in S : \forall y \in S : x \leq y$

View ID Helpers

$LeaderID(viewID) \triangleq (viewID \% Len(ReplicaOrder)) + (\text{IF } viewID \geq Len(ReplicaOrder) \text{ THEN } 1 \text{ ELSE } 0)$

$Leader(viewID) \triangleq ReplicaOrder[LeaderID(viewID)]$ **remember $\langle \rangle$ are 1-indexed**

Log Manipulation Helpers

The order of 2 *log* entries are decided by the tuple $\langle deadline, clientID, requestID \rangle$

Usually, deadline makes the two entries comparable

When 2 different entries have the same deadline, the tie is broken with *clientID*

Further, the tie is broken is *requestID*

(unnecessary if we only allow client to submit one request at one tick)

$EntryLeq(l1, l2) \triangleq \wedge l1.deadline \leq l2.deadline$
 $\wedge l1.clientID \leq l2.clientID$
 $\wedge l1.requestID \leq l2.requestID$

$EntryEq(l1, l2) \triangleq \wedge l1.deadline = l2.deadline$
 $\wedge l1.clientID = l2.clientID$
 $\wedge l1.requestID = l2.requestID$

$EntryLessThan(l1, l2) \triangleq \wedge EntryLeq(l1, l2)$
 $\wedge \neg(EntryEq(l1, l2))$

Find entry in one replica's *log* ($\langle clientID, reqID \rangle$ can uniquely identify the *log* entry)

We do not check deadline, because the leader may have modified the request's deadline

Return 0 when we fail to find it (remember Sequence is 1-indexed in TLA+, so 0 can serve as a dummy value)

$FindEntry(clientID, reqID, log) \triangleq$

LET

$entryIndexSet \triangleq \{i \in 1 \dots Len(log) : \wedge log[i].clientID = clientID$
 $\wedge log[i].reqID = reqID\}$


```

IN
  IF  $Cardinality(entryIndexSet) = 0$  THEN
    0
  ELSE
     $Pick(entryIndexSet)$ 

```

$SortLogSeq(seq) \triangleq SortSeq(seq, \text{LAMBDA } x, y : EntryLessThan(x, y))$

Given a set of logs, return the sorted *log* list

```

 $GetSortLogSeq(S) \triangleq$  LET
   $seq \triangleq Set2Seq(S)$ 
IN
   $SortLogSeq(seq)$ 

```

Merge logs, first put all *log* items together, deduplicated (*i.e.* UNION them into a set). Then, do filtering and only keep those that have appeared in at least $\lceil f/2 \rceil + 1$ replicas.

```

 $CountVotes(logll, x) \triangleq Cardinality(\{logSet \in logll : x \in logSet\})$ 

 $MergeUnSyncLogs(unSyncedLogs, lastSyncedLog) \triangleq$ 
  LET
     $unSyncedLogSet \triangleq \text{UNION } unSyncedLogs$ 
     $votedLogSet \triangleq \{x \in unSyncedLogSet :$ 
       $\wedge EntryLessThan(lastSyncedLog, x)$ 
       $\wedge CountVotes(unSyncedLogs, x) \geq RecoveryQuorumSize\}$ 
  IN
     $GetSortLogSeq(votedLogSet)$ 

```

Network Helpers

Add a message to the network

$Send(ms) \triangleq messages' = messages \cup ms$

Convert the request format to a *log* format (by summing up *s* and *l* to get deadline)

```

 $Req2Log(req) \triangleq [$ 
   $mtype \mapsto MLogEntry,$ 
   $deadline \mapsto req.s + req.l,$ 
   $clientID \mapsto req.sender,$ 
   $requestID \mapsto req.requestID,$ 
   $command \mapsto req.command$ 
 $]$ 

```

Index does not need to include command field, which is the body of the request/*log*, and can be very large

```

 $GetLogIndex(entry) \triangleq [$ 
   $mtype \mapsto MLogIndex,$ 
   $deadline \mapsto entry.deadline,$ 
   $clientID \mapsto entry.clientID,$ 
   $requestID \mapsto entry.requestID$ 
 $]$ 

```

$$\begin{aligned} \text{GetLogIndexFromReply}(\text{reply}) \triangleq & [\\ & \text{mtype} \mapsto \text{MLogIndex}, \\ & \text{deadline} \mapsto \text{reply.deadline}, \\ & \text{clientID} \mapsto \text{reply.dest}, \\ & \text{requestID} \mapsto \text{reply.requestID} \\ &] \end{aligned}$$

$$\begin{aligned} \text{IndexEq}(\text{index}, \text{msg}) \triangleq & \wedge \text{index.deadline} = \text{msg.deadline} \\ & \wedge \text{index.clientID} = \text{msg.clientID} \\ & \wedge \text{index.requestID} = \text{msg.requestID} \end{aligned}$$

Add local time to the message (for easy debug)

$$\text{Msg2RLog}(\text{msg}, r) \triangleq \text{msg} @@ [tl \mapsto v\text{ReplicaClock}[r]]$$

$$\text{LastLog}(\text{logList}) \triangleq \text{IF } \text{Len}(\text{logList}) = 0 \text{ THEN } \text{NULLLog} \text{ ELSE } \text{logList}[\text{Len}(\text{logList})]$$

$$\text{MergeCrashVector}(\text{cv1}, \text{cv2}) \triangleq [r \in \text{Replicas} \mapsto \text{Max}(\{\text{cv1}[r], \text{cv2}[r]\})]$$

$$\begin{aligned} \text{CheckCrashVector}(m, r) \triangleq & \\ & \text{IF } m.\text{cv}[m.\text{sender}] < v\text{CrashVector}[r][m.\text{sender}] \text{ THEN} \\ & \quad \text{FALSE Potential stray message} \\ & \text{ELSE} \\ & \quad v\text{CrashVector}' = [v\text{CrashVector} \text{ EXCEPT } ![r] = \text{MergeCrashVector}(m.\text{cv}, v\text{CrashVector}[r])] \end{aligned}$$

Message Handlers and Actions

Client action

Client c sends a request

We assume client can only send one request in one tick of time

If time has reached the bound, this client cannot send request any more

$$\begin{aligned} \text{ClientSendRequest}(c) \triangleq & \wedge v\text{ClientClock}[c] < \text{MaxTime} \\ & \wedge v\text{ClientReqNum}[c] < \text{MaxReqNum} \\ & \wedge \text{Send}(\{ [\text{mtype} \mapsto \text{MClientRequest}, \\ & \quad \text{sender} \mapsto c, \text{clientID} \\ & \quad \text{requestID} \mapsto v\text{ClientReqNum}[c] + 1, \text{requestID} \\ & \quad \text{command} \mapsto "", \\ & \quad s \mapsto v\text{ClientClock}[c], \text{submission time} \\ & \quad l \mapsto \text{LatencyBounds}[c], \text{latency bound} \\ & \quad \text{dest} \mapsto r \\ & \quad] : r \in \text{Replicas} \}) \\ & \wedge v\text{ClientClock}' = [v\text{ClientClock} \text{ EXCEPT } ![c] = v\text{ClientClock}[c] + 1] \\ & \wedge v\text{ClientReqNum}' = [v\text{ClientReqNum} \text{ EXCEPT } ![c] = v\text{ClientReqNum}[c] + 1] \\ & \wedge \text{UNCHANGED } \langle \text{replicaVars} \rangle \end{aligned}$$

$$\begin{aligned}
& Duplicate(entry, logSet) \triangleq \\
& \text{LET} \\
& \quad findSet \triangleq \{x \in logSet : \wedge x.clientID = entry.clientID \\
& \quad \quad \quad \wedge x.requestID = entry.requestID\} \\
& \text{IN} \\
& \quad Cardinality(findSet) > 0 \\
& \text{Replica } r \text{ receives } MClientRequest, m \\
& HandleClientRequest(r, m) \triangleq \\
& \text{LET} \\
& \quad mlog \triangleq Req2Log(m) \\
& \text{IN} \\
& \quad \text{If the request is duplicate, it will no longer be appended to the } log \\
& \quad \text{Replicas simply reply the previous execution result of this request} \\
& \quad \text{(we do not model execution in this spec)} \\
& \quad \wedge \neg Duplicate(mlog, Seq2Set(vLog[r]) \cup vEarlyBuffer[r]) \\
& \quad \wedge vReplicaStatus[r] = StNormal \\
& \quad \text{The request can enter the early buffer} \\
& \quad \wedge \vee \wedge EntryLessThan(LastLog(vLog[r]), mlog) \\
& \quad \quad \wedge vEarlyBuffer' = [\\
& \quad \quad \quad vEarlyBuffer \text{ EXCEPT } ![r] = vEarlyBuffer[r] \cup \{mlog\} \\
& \quad \quad] \\
& \quad \wedge \text{UNCHANGED } \langle networkVars, clientVars, \\
& \quad \quad vLog, vViewID, vReplicaClock, \\
& \quad \quad vLastNormView, vViewChanges, vReplicaStatus, \\
& \quad \quad vSyncPoint, vLateBuffer, \\
& \quad \quad vTentativeSync, vSyncReps, vCommitPoint, \\
& \quad \quad vUUIDCounter, vCrashVector, \\
& \quad \quad vCrashVectorReps, vRecoveryReps \rangle \\
& \quad (1) \text{ Followers' early buffers do not accept the request} \\
& \quad \quad \text{if its deadline is smaller than previously appended (last released) entry,} \\
& \quad \quad \text{so followers directly put the request into the late buffer} \\
& \quad (2) \text{ Leader modifies its deadline to be larger than the last released entry} \\
& \quad \quad \text{so as to make it eligible for entering the early buffer} \\
& \quad \vee \wedge EntryLessThan(mlog, LastLog(vLog[r])) \\
& \quad \wedge \text{IF } r = Leader(vViewID[r]) \text{ THEN } \text{this replica is the leader in the current view} \\
& \quad \quad \wedge vEarlyBuffer' = [\\
& \quad \quad \quad vEarlyBuffer \text{ EXCEPT } ![r] = vEarlyBuffer[r] \cup \{[\\
& \quad \quad \quad \quad mtype \mapsto MLogEntry, \\
& \quad \quad \quad \quad clientID \mapsto mlog.clientID, \\
& \quad \quad \quad \quad requestID \mapsto mlog.requestID, \\
& \quad \quad \quad \quad deadline \mapsto LastLog(vLog[r]).deadline + 1, \\
& \quad \quad \quad \quad command \mapsto mlog.command \\
& \quad \quad \quad]\} \\
& \quad \quad] \\
& \quad]
\end{aligned}$$

```

    ∧ UNCHANGED  ⟨networkVars, clientVars,
                  vLog, vViewID, vReplicaClock,
                  vLastNormView, vViewChanges, vReplicaStatus,
                  vSyncPoint, vLateBuffer,
                  vTentativeSync, vSyncReps, vCommitPoint,
                  vUUIDCounter, vCrashVector,
                  vCrashVectorReps, vRecoveryReps⟩

ELSE  this replica is a follower in the current view
    ∧ vLateBuffer' = [
        vLateBuffer EXCEPT ![r] = vLateBuffer[r] ∪ {mlog}
    ]

    ∧ UNCHANGED  ⟨networkVars, clientVars,
                  vLog, vEarlyBuffer, vViewID, vReplicaClock,
                  vLastNormView, vViewChanges, vReplicaStatus,
                  vSyncPoint, vTentativeSync,
                  vSyncReps, vCommitPoint,
                  vUUIDCounter, vCrashVector,
                  vCrashVectorReps, vRecoveryReps⟩

Release relevant requests from vEarlyBuffer and append to vLog,
and then send a fast reply
FlushEarlyBuffer(r) ≜
    LET
        validLogSet ≜ {x ∈ vEarlyBuffer[r] :
            ∧ x.deadline < vReplicaClock[r] < rather than ≤
            ∧ EntryLessThan(LastLog(vLog[r]), x)}
        validLogs ≜ GetSortLogSeq(validLogSet)
        newLogStart ≜ Len(vLog[r]) + 1
    IN
        ∧ vLog' = [vLog EXCEPT ![r] = vLog[r] ∘ validLogs]
        ∧ vEarlyBuffer' = [vEarlyBuffer EXCEPT ![r]
            = {x ∈ vEarlyBuffer[r] : x.deadline ≥ vReplicaClock[r]}] ≥ rather than >
        ∧ Send({[mtype ↦ MFastReply,
            sender ↦ r,
            dest ↦ vLog'[r][i].clientID,
            viewID ↦ vViewID[r],
            requestID ↦ vLog'[r][i].requestID,
            hash ↦ [
                log ↦ SubSeq(vLog'[r], 1, i),
                cv ↦ vCrashVector
            ],
            deadline ↦ vLog'[r][i].deadline,
            logSlotNum ↦ i
        ] : i ∈ newLogStart .. Len(vLog'[r])})

```

```

 $\wedge$  IF  $r = \text{Leader}(v\text{ViewID}[r])$  THEN
   $\wedge v\text{SyncPoint}' = [v\text{SyncPoint} \text{ EXCEPT } ![r] = \text{Len}(v\text{Log}'[r])]$ 
   $\wedge$  UNCHANGED  $\langle \text{clientVars}, v\text{ViewID}, v\text{LastNormView}, v\text{ViewChanges},$ 
     $v\text{ReplicaStatus}, v\text{ReplicaClock}, v\text{LateBuffer},$ 
     $v\text{TentativeSync}, v\text{SyncReps}, v\text{CommitPoint},$ 
     $v\text{UUIDCounter}, v\text{CrashVector},$ 
     $v\text{CrashVectorReps}, v\text{RecoveryReps} \rangle$ 
ELSE
  UNCHANGED  $\langle \text{clientVars}, v\text{ViewID}, v\text{LastNormView}, v\text{ViewChanges},$ 
     $v\text{ReplicaStatus}, v\text{ReplicaClock},$ 
     $v\text{SyncPoint}, v\text{LateBuffer},$ 
     $v\text{TentativeSync}, v\text{SyncReps}, v\text{CommitPoint},$ 
     $v\text{UUIDCounter}, v\text{CrashVector},$ 
     $v\text{CrashVectorReps}, v\text{RecoveryReps} \rangle$ 

```

Clock can be random value ($\text{RandomElement}(1 \dots \text{MaxTime})$),

because clock sync algorithm can give negative offset, or even fails

But *Nezha* depend on clock for performance but not for correctness

If the replica clock goes beyond MaxTime , it will stop processing

Since Clock is moved, then replicas can release relevant requests and append to logs

```

 $\text{ReplicaClockMove}(r) \triangleq \wedge$  IF  $v\text{ReplicaClock}[r] < \text{MaxTime}$  THEN
   $v\text{ReplicaClock}' = [$ 
     $v\text{ReplicaClock} \text{ EXCEPT } ![r] = \text{RandomElement}(1 \dots \text{MaxTime})$ 
   $]$ 
ELSE
  UNCHANGED  $v\text{ReplicaClock}$ 
 $\wedge$  UNCHANGED  $\langle \text{networkVars}, \text{clientVars},$ 
   $v\text{Log}, v\text{EarlyBuffer}, v\text{ViewID},$ 
   $v\text{LastNormView}, v\text{ViewChanges}, v\text{ReplicaStatus},$ 
   $v\text{SyncPoint}, v\text{LateBuffer}, v\text{TentativeSync},$ 
   $v\text{SyncReps}, v\text{CommitPoint},$ 
   $v\text{UUIDCounter}, v\text{CrashVector},$ 
   $v\text{CrashVectorReps}, v\text{RecoveryReps} \rangle$ 

```

Client clock move does not change any other things

```

 $\text{ClientClockMove}(c) \triangleq \wedge$  IF  $v\text{ClientClock}[c] < \text{MaxTime}$  THEN
   $v\text{ClientClock}' = [$ 
     $v\text{ClientClock} \text{ EXCEPT } ![c] = \text{RandomElement}(1 \dots \text{MaxTime})$ 
   $]$ 
ELSE
  UNCHANGED  $v\text{ClientClock}$ 
 $\wedge$  UNCHANGED  $\langle \text{networkVars}, \text{replicaVars}, v\text{ClientReqNum} \rangle$ 

```

Index Synchronization to Fix Set Inequality

Leader replica r starts index synchronization

$$\begin{aligned}
& \text{StartIndexSync}(r) \triangleq \\
& \text{LET} \\
& \quad \text{indices} \triangleq \{ \text{GetLogIndex}(\text{vLog}[r][i]) : i \in 1 \dots \text{Len}(\text{vLog}[r]) \} \\
& \text{IN} \\
& \quad \wedge r = \text{Leader}(\text{vViewID}[r]) \\
& \quad \wedge \text{vReplicaStatus}[r] = \text{StNormal} \\
& \quad \wedge \text{Cardinality}(\text{indices}) > 0 \quad \text{leader has log entries to sync} \\
& \quad \wedge \text{Send}(\{ \text{mtype} \mapsto \text{MIndexSync}, \\
& \quad \quad \text{sender} \mapsto r, \\
& \quad \quad \text{dest} \mapsto d, \\
& \quad \quad \text{viewID} \mapsto \text{vViewID}[r], \\
& \quad \quad \text{logindcies} \mapsto \text{indices} \} : d \in \text{Replicas} \}) \\
& \quad \wedge \text{UNCHANGED} \langle \text{clientVars}, \text{replicaVars} \rangle
\end{aligned}$$

$$\begin{aligned}
& \text{GetSyncLogs}(\text{logSeq}, \text{indices}) \triangleq \\
& \quad \text{LET} \\
& \quad \quad \text{logSet} \triangleq \{ l \in \text{Seq2Set}(\text{logSeq}) : \exists \text{index} \in \text{indices} : \text{EntryEq}(\text{index}, l) \} \\
& \quad \text{IN} \\
& \quad \quad \text{GetSortLogSeq}(\text{logSet})
\end{aligned}$$

$$\begin{aligned}
& \text{GetUnSyncLogs}(\text{logSeq}, \text{lastSyncedLog}) \triangleq \\
& \quad \text{LET} \\
& \quad \quad \text{logSet} \triangleq \{ l \in \text{Seq2Set}(\text{logSeq}) : \text{EntryLessThan}(\text{lastSyncedLog}, l) \} \\
& \quad \text{IN} \\
& \quad \quad \text{GetSortLogSeq}(\text{logSet})
\end{aligned}$$

Replica r receives *IndexSync* message, m

$$\begin{aligned}
& \text{HandleIndexSync}(r, m) \triangleq \\
& \quad \wedge r \neq \text{Leader}(\text{vViewID}[r]) \\
& \quad \wedge \text{vReplicaStatus}[r] = \text{StNormal} \\
& \quad \wedge m.\text{viewID} = \text{vViewID}[r] \\
& \quad \wedge m.\text{sender} = \text{Leader}(\text{vViewID}[r]) \\
& \quad \wedge \text{vSyncPoint}[r] < \text{Len}(m.\text{logindcies}) \\
& \quad \wedge \text{LET} \\
& \quad \quad \text{entries} \triangleq \{ \text{vLog}[r][i] : i \in 1 \dots \text{Len}(\text{vLog}[r]) \} \\
& \quad \quad \text{indices} \triangleq \{ \text{GetLogIndex}(\text{vLog}[r][i]) : i \in 1 \dots \text{Len}(\text{vLog}[r]) \} \\
& \quad \quad \text{missedEntries} \triangleq m.\text{indices} \setminus \text{indices} \\
& \quad \text{IN} \\
& \quad \quad \text{Missing some log entries} \rightarrow \text{Send } \text{MMissEntryRequest} \\
& \quad \text{IF } \text{Cardinality}(\text{missedEntries}) > 0 \text{ THEN} \\
& \quad \quad \wedge \text{Send}(\{ \text{mtype} \mapsto \text{MMissEntryRequest}, \\
& \quad \quad \quad \text{sender} \mapsto r, \\
& \quad \quad \quad \text{dest} \mapsto d, \\
& \quad \quad \quad \text{viewID} \mapsto \text{vViewID}[r],
\end{aligned}$$

$miss \mapsto missedEntries] : d \in (Replicas \setminus \{r\})\}$

\wedge UNCHANGED $\langle vLog, vSyncPoint \rangle$

No missing entries, update $vLog$ and $vSyncPoint$, and send relevant slow replies

ELSE

LET

$syncLogs \triangleq GetSyncLogs(vLog[r], indices)$

$unsyncLogs \triangleq GetUnSyncLogs(vLog[r], LastLog(syncLogs))$

IN

$\wedge vLog' = [vLog \text{ EXCEPT } ![r] = syncLogs \circ unsyncLogs]$

$\wedge vSyncPoint' = [vSyncPoint \text{ EXCEPT } ![r] = Len(syncLogs)]$

$\wedge Send(\{[\begin{array}{ll} mtype & \mapsto MSlowReply, \\ sender & \mapsto r, \\ dest & \mapsto vLog'[r][i].clientID, \\ viewID & \mapsto vViewID[r], \\ requestID & \mapsto vLog'[r][i].requestID, \\ logSlotNum & \mapsto i \end{array}] : i \in (1 \dots Len(syncLogs))\})$

\wedge UNCHANGED $\langle clientVars, vEarlyBuffer, vViewID, vReplicaClock, vLastNormView, vViewChanges, vReplicaStatus, vLateBuffer, vTentativeSync, vSyncReps, vCommitPoint, vUUIDCounter, vCrashVector, vCrashVectorReps, vRecoveryReps \rangle$

$FindEntries(log, indices) \triangleq$
 $\{l \in Seq2Set(log) : \exists x \in indices : IndexEq(l, x)\}$

Replica r receives a request from other replicas, asking for a missing log entry

$HandleMissEntryRequest(r, m) \triangleq$

$\wedge m.viewID = vViewID[r]$

\wedge LET

$findentries \triangleq FindEntries(vLog[r], m.miss)$

IN

$\wedge Cardinality(findentries) > 0$

$\wedge Send(\{[\begin{array}{ll} mtype & \mapsto MMissEntryReply, \\ sender & \mapsto r, \\ dest & \mapsto m.sender, \\ viewID & \mapsto vViewID[r], \\ entries & \mapsto findentries \end{array}]\})$

\wedge UNCHANGED $\langle clientVars, replicaVars \rangle$

Replica r receives a reply from other replicas, providing the missing entries

$HandleMissEntryReply(r, m) \triangleq$

$\wedge m.viewID = vViewID[r]$

\wedge LET

$$\begin{aligned}
& mergedSet \triangleq Seq2Set(vLog[r]) \cup m.entries \\
& \text{IN} \\
& vLog' = [vLog \text{ EXCEPT } ![r] = GetSortLogSeq(mergedSet)] \\
& \wedge \text{UNCHANGED } \langle networkVars, clientVars, \\
& \quad vEarlyBuffer, vViewID, vReplicaClock, \\
& \quad vLastNormView, vViewChanges, vReplicaStatus, \\
& \quad vSyncPoint, vLateBuffer, \\
& \quad vTentativeSync, vSyncReps, vCommitPoint, \\
& \quad vUUIDCounter, vCrashVector, \\
& \quad vCrashVectorReps, vRecoveryReps \rangle
\end{aligned}$$

Replica Rejoin

Failed replica loses all states

$$\begin{aligned}
& StartReplicaFail(r) \triangleq \\
& \quad \wedge NumofReplicas(StRecovering) < F \quad \text{We assume at most } F \text{ replicas can fail at the same time} \\
& \quad \wedge vReplicaStatus' = [vReplicaStatus \text{ EXCEPT } ![r] = StRecovering] \\
& \quad \wedge vLog' = [vLog \text{ EXCEPT } ![r] = \langle \rangle] \\
& \quad \wedge vEarlyBuffer' = [vEarlyBuffer \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge vViewID' = [vViewID \text{ EXCEPT } ![r] = 1] \\
& \quad \wedge vLastNormView' = [vLastNormView \text{ EXCEPT } ![r] = 1] \\
& \quad \wedge vViewChanges' = [vViewChanges \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge vSyncPoint' = [vSyncPoint \text{ EXCEPT } ![r] = 0] \\
& \quad \wedge vLateBuffer' = [vLateBuffer \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge vTentativeSync' = [vTentativeSync \text{ EXCEPT } ![r] = 0] \\
& \quad \wedge vSyncReps' = [vSyncReps \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge vCommitPoint' = [vCommitPoint \text{ EXCEPT } ![r] = 0] \\
& \quad \wedge vCrashVector' = [vCrashVector \text{ EXCEPT } ![r] = [rr \in Replicas \mapsto 0]] \\
& \quad \wedge vCrashVectorReps' = [vCrashVectorReps \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge vRecoveryReps' = [vRecoveryReps \text{ EXCEPT } ![r] = \{\}] \\
& \quad \wedge \text{UNCHANGED } \langle vReplicaClock, vUUIDCounter, clientVars, networkVars \rangle
\end{aligned}$$

Recovering replica starts recovery (by first sending *CrashVectorReq*)

$$\begin{aligned}
& StartReplicaRecovery(r) \triangleq \\
& \quad \wedge vReplicaStatus[r] = StRecovering \\
& \quad \wedge vUUIDCounter' = [vUUIDCounter \text{ EXCEPT } ![r] = vUUIDCounter[r] + 1] \\
& \quad \wedge Send(\{[mtype \mapsto MCrashVectorReq, \\
& \quad \quad sender \mapsto r, \\
& \quad \quad dest \mapsto d, \\
& \quad \quad nonce \mapsto vUUIDCounter'[r]] : d \in Replicas\}) \\
& \quad \wedge \text{UNCHANGED } \langle vLog, vEarlyBuffer, vViewID, vReplicaClock, \\
& \quad \quad vLastNormView, vViewChanges, vReplicaStatus,
\end{aligned}$$

$vSyncPoint, vLateBuffer,$
 $vTentativeSync, vSyncReps, vCommitPoint,$
 $vCrashVector, vCrashVectorReps, vRecoveryReps,$
 $clientVars \rangle$

$HandleCrashVectorReq(r, m) \triangleq$
 $\wedge vReplicaStatus[r] = StNormal$
 $\wedge Send(\{[mtype \mapsto MCrashVectorRep,$
 $sender \mapsto r,$
 $dest \mapsto m.sender,$
 $nonce \mapsto m.nonce,$
 $cv \mapsto vCrashVector[r]]\})$
 $\wedge UNCHANGED \langle replicaVars, clientVars \rangle$

$HandleCrashVectorRep(r, m) \triangleq$
 $\wedge vReplicaStatus[r] = StRecovering$
 $\wedge vUUIDCounter[r] = m.nonce$
 $\wedge Cardinality(vCrashVectorReps[r]) \leq F$
 $\wedge \neg DuplicateRep(vCrashVectorReps[r], m)$
 $\wedge vCrashVectorReps' = [vCrashVectorReps \text{ EXCEPT } ![r] = vCrashVectorReps[r] \cup \{m\}]$
 $\wedge vCrashVector' = [vCrashVector \text{ EXCEPT } ![r] = MergeCrashVector(vCrashVector[r], m.cv)]$
 $\wedge \text{IF } Cardinality(vCrashVectorReps') = F + 1 \text{ THEN } \text{got enough replies and can settle down } cv$
 $Send(\{[mtype \mapsto MRecoveryReq,$
 $sender \mapsto r,$
 $dest \mapsto d,$
 $nonce \mapsto m.nonce,$
 $cv \mapsto vCrashVector'[r] : d \in Replicas\})$
 $ELSE$
 $UNCHANGED \langle networkVars \rangle$
 $\wedge UNCHANGED \langle vLog, vEarlyBuffer, vViewID, vReplicaClock,$
 $vLastNormView, vViewChanges, vReplicaStatus,$
 $vSyncPoint, vLateBuffer,$
 $vTentativeSync, vSyncReps, vCommitPoint,$
 $vUUIDCounter, vRecoveryReps,$
 $clientVars \rangle$

$HandleRecoveryReq(r, m) \triangleq$
 $\wedge vReplicaStatus[r] = StNormal$
 $\wedge vCrashVector' = [vCrashVector \text{ EXCEPT } ![r] = MergeCrashVector(vCrashVector[r], m.cv)]$
 $\wedge Send(\{[mtype \mapsto MRecoveryRep,$
 $sender \mapsto r,$
 $dest \mapsto m.sender,$

$$\begin{aligned}
& viewID \mapsto vViewID[r], \\
& cv \mapsto vCrashVector'[r] : d \in Replicas \} \\
& \wedge \text{UNCHANGED } \langle vLog, vEarlyBuffer, vViewID, vReplicaClock, \\
& \quad vLastNormView, vViewChanges, vReplicaStatus, \\
& \quad vSyncPoint, vLateBuffer, \\
& \quad vTentativeSync, vSyncReps, vCommitPoint, \\
& \quad vUUIDCounter, vCrashVectorReps, vRecoveryReps, \\
& \quad clientVars \rangle
\end{aligned}$$

$$\begin{aligned}
& HandleRecoveryRep(r, m) \triangleq \\
& \quad \wedge vReplicaStatus[r] = StRecovering \\
& \quad \wedge Cardinality(vRecoveryReps[r]) \leq F \\
& \quad \wedge \neg DuplicateRep(vRecoveryReps[r], m.sender) \\
& \quad \wedge CheckCrashVector(m, r) \\
& \quad \wedge vRecoveryReps' = [vRecoveryReps \text{ EXCEPT } ![r] = vRecoveryReps[r] \cup \{m\}] \\
& \quad \wedge \text{IF } Cardinality(vRecoveryReps') = F + 1 \text{ THEN } \text{got enough replies} \\
& \quad \quad \text{LET} \\
& \quad \quad \quad newView \triangleq Max(\{mm.viewID : mm \in vRecoveryReps'[r]\}) \\
& \quad \quad \quad leaderId \triangleq newView \% Cardinality(Replicas) \\
& \quad \quad \text{IN} \\
& \quad \quad \quad Send(\{[mtype \mapsto MStateTransferReq, \\
& \quad \quad \quad \quad sender \mapsto r, \\
& \quad \quad \quad \quad dest \mapsto leaderId, \\
& \quad \quad \quad \quad cv \mapsto vCrashVector'[r] : d \in Replicas\}) \\
& \quad \quad \text{ELSE} \\
& \quad \quad \quad \text{UNCHANGED } \langle networkVars \rangle \\
& \quad \wedge \text{UNCHANGED } \langle vLog, vEarlyBuffer, vViewID, vReplicaClock, \\
& \quad \quad vLastNormView, vViewChanges, vReplicaStatus, \\
& \quad \quad vSyncPoint, vLateBuffer, \\
& \quad \quad vTentativeSync, vSyncReps, vCommitPoint, \\
& \quad \quad vUUIDCounter, vCrashVectorReps, \\
& \quad \quad clientVars \rangle
\end{aligned}$$

$$\begin{aligned}
& HandleStateTransferReq(r, m) \triangleq \\
& \quad \wedge vReplicaStatus[r] = StNormal \\
& \quad \wedge CheckCrashVector(m, r) \\
& \quad \wedge Send(\{[mtype \mapsto MStateTransferRep, \\
& \quad \quad \quad sender \mapsto r, \\
& \quad \quad \quad dest \mapsto m.sender, \\
& \quad \quad \quad log \mapsto vLog[r], \\
& \quad \quad \quad sp \mapsto vSyncPoint[r],
\end{aligned}$$

$$\begin{aligned}
cp &\mapsto vCommitPoint[r], \\
cv &\mapsto vCrashVector'[r]\} \\
\wedge \text{UNCHANGED } &\langle vLog, vEarlyBuffer, vViewID, vReplicaClock, \\
&vLastNormView, vViewChanges, vReplicaStatus, \\
&vSyncPoint, vLateBuffer, \\
&vTentativeSync, vSyncReps, vCommitPoint, \\
&vUUIDCounter, vCrashVectorReps, vRecoveryReps, \\
&clientVars \rangle
\end{aligned}$$

$$\begin{aligned}
HandleStateTransferRep(r, m) &\triangleq \\
&\wedge vReplicaStatus[r] = StRecovering \\
&\wedge CheckCrashVector(m, r) \\
&\wedge vLog' = [vLog \text{ EXCEPT } ![r] = m.log] \\
&\wedge vSyncPoint' = [vSyncPoint \text{ EXCEPT } ![r] = m.sp] \\
&\wedge vCommitPoint' = [vCommitPoint \text{ EXCEPT } ![r] = m.cp] \\
&\wedge vViewID' = [vViewID \text{ EXCEPT } ![r] = m.viewID] \\
&\wedge vEarlyBuffer' = [vEarlyBuffer \text{ EXCEPT } ![r] = \{\}] \\
&\wedge vLastNormView' = [vLastNormView \text{ EXCEPT } ![r] = m.viewID] \\
&\wedge vViewChanges' = [vViewChanges \text{ EXCEPT } ![r] = \{\}] \\
&\wedge vReplicaStatus' = [vReplicaStatus \text{ EXCEPT } ![r] = StNormal] \\
&\wedge vLateBuffer' = [vLateBuffer \text{ EXCEPT } ![r] = \{\}] \\
&\wedge vTentativeSync' = [vTentativeSync \text{ EXCEPT } ![r] = m.sp] \\
&\wedge vSyncReps' = [vSyncReps \text{ EXCEPT } ![r] = \{\}] \\
&\wedge vCrashVectorReps' = [vCrashVectorReps \text{ EXCEPT } ![r] = \{\}] \\
&\wedge vRecoveryReps' = [vRecoveryReps \text{ EXCEPT } ![r] = \{\}] \\
&\wedge \text{UNCHANGED } \langle vReplicaClock, vUUIDCounter, clientVars \rangle
\end{aligned}$$

Leader Change

Replica r starts a *Leader change*

$$\begin{aligned}
StartLeaderChange(r) &\triangleq \\
&\wedge Send(\{[mtype \mapsto MViewChangeReq, \\
&\quad sender \mapsto r, \\
&\quad dest \mapsto d, \\
&\quad viewID \mapsto vViewID[r] + 1, \\
&\quad cv \mapsto vCrashVector[r]] : d \in Replicas\}) \\
&\wedge \text{UNCHANGED } \langle replicaVars, clientVars \rangle
\end{aligned}$$

View Change Handlers

Replica r gets *MViewChangeReq*, m

$$\begin{aligned}
HandleViewChangeReq(r, m) &\triangleq \\
&\text{LET}
\end{aligned}$$

$$\begin{aligned}
currentViewID &\triangleq vViewID[r] \\
newViewID &\triangleq Max(\{currentViewID, m.viewID\}) \\
newLeaderNum &\triangleq LeaderID(newViewID)
\end{aligned}$$

IN

Recovering replica does not participate in view change

$$\begin{aligned}
&\wedge vReplicaStatus[r] \neq StRecovering \\
&\wedge currentViewID \neq newViewID \\
&\wedge CheckCrashVector(m, r) \\
&\wedge vReplicaStatus' = [vReplicaStatus \text{ EXCEPT } ![r] = StViewChange] \\
&\wedge vViewID' = [vViewID \text{ EXCEPT } ![r] = newViewID] \\
&\wedge vViewChanges' = [vViewChanges \text{ EXCEPT } ![r] = \{\}] \\
&\wedge Send(\{[mtype \mapsto MViewChange, \\
&\quad dest \mapsto Leader(newViewID), \\
&\quad sender \mapsto r, \\
&\quad viewID \mapsto newViewID, \\
&\quad lastNormal \mapsto vLastNormView[r], \\
&\quad syncedLog \mapsto SubSeq(vLog[r], 1, vSyncPoint[r]), \\
&\quad unsyncedLog \mapsto SubSeq(vLog[r], vSyncPoint[r] + 1, Len(vLog[r])), \\
&\quad cv \mapsto vCrashVector[r]]\} \cup \\
&\quad \text{Send the } MViewChangeReqs \text{ in case this is an entirely new view} \\
&\quad \{[mtype \mapsto MViewChangeReq, \\
&\quad sender \mapsto r, \\
&\quad dest \mapsto d, \\
&\quad viewID \mapsto newViewID, \\
&\quad cv \mapsto vCrashVector[r]] : d \in Replicas\}) \\
&\wedge UNCHANGED \langle clientVars, vLog, vEarlyBuffer, vReplicaClock, \\
&\quad vLastNormView, vSyncPoint, vLateBuffer, \\
&\quad vTentativeSync, vSyncReps, vCommitPoint, \\
&\quad vUUIDCounter, vCrashVectorReps, vRecoveryReps \rangle
\end{aligned}$$

Replica r receives $MViewChange, m$

$$HandleViewChange(r, m) \triangleq$$

Recovering replica does not participate in view change

$$\begin{aligned}
&\wedge vReplicaStatus[r] \neq StRecovering \\
&\quad \text{Add the message to the log} \\
&\wedge vViewID[r] = m.viewID \\
&\wedge vReplicaStatus[r] = StViewChange \\
&\quad \text{This replica is the leader} \\
&\wedge Leader(vViewID[r]) = r \\
&\wedge CheckCrashVector(m, r) \\
&\wedge vViewChanges' = [vViewChanges \text{ EXCEPT } ![r] = vViewChanges[r] \cup \{m\}] \\
&\quad \text{If there's enough replies, start the new view} \\
&\wedge LET \\
&\quad isViewPromise(M) \triangleq \wedge \{n.sender : n \in M\} \in Quorums
\end{aligned}$$

$$vCMs \triangleq \{n \in M : n.sender = r \wedge n.mtype = MViewChange \wedge n.viewID = vViewID[r]\}$$

Create the state for the new view

$$normalViews \triangleq \{n.lastNormal : n \in vCMs\}$$

Choose the largest normal view (i.e. the newest)

$$lastNormal \triangleq (\text{CHOOSE } v \in normalViews : \forall v2 \in normalViews : v2 \leq v)$$

For logs before $vSyncPoint$ (i.e. $syncedLog$), we directly copy from the $bestCandidates$
 For $unsyncedLog$, we do quorum check to decide which ones should be added to recovery Log

$$goodCandidates \triangleq \{o \in vCMs : o.lastNormal = lastNormal\}$$

$bestCandidate$ can only be picked from $goodCandidates$,
 because previous views may include invalid logs

$$bestCandidate \triangleq \text{CHOOSE } n \in goodCandidates : \forall y \in goodCandidates : Len(n.syncedLog) \geq Len(y.syncedLog)$$

$$unSyncedLogs \triangleq \{Seq2Set(n.unsyncedLog) : n \in goodCandidates\}$$

IN
 IF $isViewPromise(vCMs)$ THEN

$$Send(\{[mtype \mapsto MStartView, dest \mapsto d, viewID \mapsto vViewID[r], log \mapsto bestCandidate.syncedLog \circ MergeUnSyncLogs(unSyncedLogs, LastLog(bestCandidate.syncedLog))]$$

$$: d \in Replicas\})$$

 ELSE
 UNCHANGED $networkVars$

$$\wedge \text{UNCHANGED } \langle clientVars, vLog, vEarlyBuffer, vViewID, vReplicaClock, vLastNormView, vReplicaStatus, vSyncPoint, vLateBuffer, vTentativeSync, vSyncReps, vCommitPoint, vUUIDCounter, vCrashVectorReps, vRecoveryReps \rangle$$

Replica r receives a $MStartView, m$

$$HandleStartView(r, m) \triangleq$$

$$\wedge vReplicaStatus[r] \neq StRecovering$$

$$\wedge \vee vViewID[r] < m.viewID$$

$$\vee vViewID[r] = m.viewID \wedge vReplicaStatus[r] = StViewChange$$

$$\wedge CheckCrashVector(m, r)$$

$$\wedge vLog' = [vLog \text{ EXCEPT } ![r] = m.log]$$

$$\wedge vReplicaStatus' = [vReplicaStatus \text{ EXCEPT } ![r] = StNormal]$$

$$\wedge vViewID' = [vViewID \text{ EXCEPT } ![r] = m.viewID]$$

$$\wedge vLastNormView' = [vLastNormView \text{ EXCEPT } ![r] = m.viewID]$$

$$\wedge vEarlyBuffer' = [vEarlyBuffer \text{ EXCEPT } ![r] = \{\}] \text{ clear Early Buffer for the new view}$$

$$\wedge vLateBuffer' = [vLateBuffer \text{ EXCEPT } ![r] = \{\}] \text{ clear Late Buffer for the new view}$$

$\wedge vSyncPoint' = [vSyncPoint \text{ EXCEPT } ![r] = Len(m.log)]$
 $\wedge vTentativeSync' = [vTentativeSync \text{ EXCEPT } ![r] = Len(m.log)]$
 Send replies (in the new view) for all *log* items
 $\wedge \text{IF } r = Leader(m.viewID) \text{ THEN } \quad \text{Leader only sends fast reply}$
 $Send(\{[\begin{array}{ll} mtype & \mapsto MFastReply, \\ sender & \mapsto r, \\ dest & \mapsto m.log[i].clientID, \\ viewID & \mapsto m.viewID, \\ requestID & \mapsto m.log[i].requestID, \\ hash & \mapsto [\\ & \quad log \mapsto SubSeq(m.log, 1, i), \\ & \quad cv \mapsto vCrashVector \\ &] , \\ deadline & \mapsto m.log[i].deadline, \\ logSlotNum \mapsto i] : i \in (1 \dots Len(m.log)) \} \})$
 ELSE While staring view, followers knows the *log* is synced with the leader, so send slow-reply
 $Send(\{[\begin{array}{ll} mtype & \mapsto MSlowReply, \\ sender & \mapsto r, \\ dest & \mapsto m.log[i].clientID, \\ viewID & \mapsto m.viewID, \\ requestID & \mapsto m.log[i].requestID, \\ logSlotNum \mapsto i] : i \in (1 \dots Len(m.log)) \} \})$
 $\wedge \text{UNCHANGED } \langle clientVars, vReplicaClock, vViewChanges, \\ vSyncReps, vCommitPoint, vCrashVector, \\ vUUIDCounter, vCrashVectorReps, vRecoveryReps \rangle$

Periodic Synchronization

Leader replica *r* conduct synchronization periodically

This periodic sync process is different from index sync process

It ensures that all replicas' logs are stable up to their *CommitPoint* (for fast recovery)

Our *CommitPoint* is essentially the *sync-point* defined in *NOPaxos* paper

Just as mentioned in *NOPaxos* paper, it is an optional optimization for fast recovery

Nezha still works even without this part

$StartSync(r) \triangleq$
 $\wedge Leader(vViewID[r]) = r$
 $\wedge vReplicaStatus[r] = StNormal$
 $\wedge vTentativeSync[r] < Len(vLog[r]) \quad \text{If } \geq \text{ then no need to sync}$
 $\wedge vSyncReps' = [vSyncReps \text{ EXCEPT } ![r] = \{\}]$
 $\wedge vTentativeSync' = [vTentativeSync \text{ EXCEPT } ![r] = Len(vLog[r])]$
 $\wedge Send(\{[\begin{array}{ll} mtype & \mapsto MSyncPrepare, \\ sender & \mapsto r, \\ dest & \mapsto d, \\ viewID & \mapsto vViewID[r], \\ log & \mapsto vLog[r] : d \in Replicas \} \})$

\wedge UNCHANGED $\langle clientVars, vLog, vEarlyBuffer, vViewID, vReplicaClock,$
 $vLastNormView, vViewChanges, vReplicaStatus,$
 $vSyncPoint, vLateBuffer, vCommitPoint,$
 $vUUIDCounter, vCrashVector,$
 $vCrashVectorReps, vRecoveryReps \rangle$

Replica r receives $MSyncPrepare, m$
 $HandleSyncPrepare(r, m) \triangleq$
 LET
 $newLog \triangleq m.log \circ GetUnSyncLogs(vLog[r], LastLog(m.log))$
 IN
 $\wedge vReplicaStatus[r] = StNormal$
 $\wedge m.viewID = vViewID[r]$
 $\wedge m.sender = Leader(vViewID[r])$
 \wedge IF $vSyncPoint[r] < Len(m.log)$ THEN
 $\wedge vSyncPoint' = [vSyncPoint \text{ EXCEPT } ![r] = Len(m.log)]$
 $\wedge vLog' = [vLog \text{ EXCEPT } ![r] = newLog]$
 $\wedge Send(\{[$ $mtype \mapsto MSlowReply,$
 $sender \mapsto r,$
 $dest \mapsto m.log[i].clientID,$
 $viewID \mapsto m.viewID,$
 $requestID \mapsto m.log[i].requestID,$
 $logSlotNum \mapsto i] : i \in (1 \dots Len(m.log))\})$
 ELSE
 UNCHANGED $\langle vLog, vSyncPoint \rangle$
 $\wedge Send(\{[$ $mtype \mapsto MSyncRep,$
 $sender \mapsto r,$
 $dest \mapsto m.sender,$
 $viewID \mapsto vViewID[r],$
 $logSlotNumber \mapsto Len(m.log)]\}$
 $)$
 \wedge UNCHANGED $\langle clientVars, vEarlyBuffer, vViewID, vReplicaClock,$
 $vLastNormView, vViewChanges, vReplicaStatus,$
 $vLateBuffer, vTentativeSync, vSyncReps, vCommitPoint,$
 $vUUIDCounter, vCrashVector,$
 $vCrashVectorReps, vRecoveryReps \rangle$

Replica r receives $MSyncRep, m$
 $HandleSyncRep(r, m) \triangleq$
 $\wedge m.viewID = vViewID[r]$
 $\wedge vReplicaStatus[r] = StNormal$
 $\wedge vSyncReps' = [vSyncReps \text{ EXCEPT } ![r] = vSyncReps[r] \cup \{m\}]$
 \wedge LET $isViewPromise(M) \triangleq \wedge \{n.sender : n \in M\} \in Quorums$

$$\begin{aligned}
sRMs & \triangleq \wedge \exists n \in M : n.sender = r \\
& \triangleq \{n \in vSyncReps'[r] : \\
& \quad \wedge n.mtype = MSyncRep \\
& \quad \wedge n.viewID = vViewID[r] \\
& \quad \wedge n.logSlotNumber = vTentativeSync[r]\} \\
committedLog & \triangleq \text{IF } vTentativeSync[r] \geq 1 \text{ THEN} \\
& \quad SubSeq(vLog[r], 1, vTentativeSync[r]) \\
& \quad \text{ELSE} \\
& \quad \langle \rangle \\
\text{IN} \\
& \text{IF } isViewPromise(sRMs) \text{ THEN} \\
& \quad \wedge Send(\{[mtype \mapsto MSyncCommit, \\
& \quad \quad sender \mapsto r, \\
& \quad \quad dest \mapsto d, \\
& \quad \quad viewID \mapsto vViewID[r], \\
& \quad \quad log \mapsto committedLog] : \\
& \quad \quad d \in Replicas\}) \\
& \quad \wedge vCommitPoint' = [vCommitPoint \text{ EXCEPT } ![r] = vTentativeSync[r]] \\
& \quad \text{ELSE} \\
& \quad \text{UNCHANGED } \langle networkVars, vCommitPoint \rangle \\
& \wedge \text{UNCHANGED } \langle clientVars, vLog, vEarlyBuffer, vViewID, \\
& \quad vReplicaClock, vLastNormView, vViewChanges, \\
& \quad vReplicaStatus, vSyncPoint, vLateBuffer, \\
& \quad vTentativeSync, vUUIDCounter, vCrashVector, \\
& \quad vCrashVectorReps, vRecoveryReps \rangle
\end{aligned}$$

Replica r receives $MSyncCommit, m$

$$\begin{aligned}
HandleSyncCommit(r, m) & \triangleq \\
& \text{LET} \\
& \quad newLog \triangleq m.log \circ GetUnSyncLogs(vLog[r], LastLog(m.log)) \\
& \text{IN} \\
& \quad \wedge vReplicaStatus[r] = StNormal \\
& \quad \wedge m.viewID = vViewID[r] \\
& \quad \wedge m.sender = Leader(vViewID[r]) \\
& \quad \wedge \text{IF } Len(m.log) \leq vCommitPoint[r] \text{ THEN} \\
& \quad \quad \text{UNCHANGED } \langle vCommitPoint, vLog \rangle \\
& \quad \text{ELSE} \\
& \quad \quad \wedge vLog' = [vLog \text{ EXCEPT } ![r] = newLog] \\
& \quad \quad \wedge vCommitPoint' = [vCommitPoint \text{ EXCEPT } ![r] = Len(m.log)] \\
& \quad \quad \wedge Send(\{[mtype \mapsto MSlowReply, \\
& \quad \quad \quad sender \mapsto r, \\
& \quad \quad \quad dest \mapsto m.log[i].clientID, \\
& \quad \quad \quad viewID \mapsto m.viewID, \\
& \quad \quad \quad requestID \mapsto m.log[i].requestID,
\end{aligned}$$

$$\wedge \text{UNCHANGED } \langle \text{networkVars}, \text{clientVars}, \text{vEarlyBuffer}, \\ \text{vViewID}, \text{vReplicaClock}, \text{vLastNormView}, \text{vViewChanges}, \\ \text{vReplicaStatus}, \text{vSyncPoint}, \text{vLateBuffer}, \\ \text{vTentativeSync}, \text{vSyncReps}, \\ \text{vUUIDCounter}, \text{vCrashVector}, \\ \text{vCrashVectorReps}, \text{vRecoveryReps} \rangle$$

Invariants and Helper Functions

A request/log is committed in two possible cases:

- (1) A fast quorum has sent either slow-reply messages, or fast-reply messages with consistent hashes [Fast Path]
- (2) A simple quorum has sent slow-reply messages [Slow Path] Both quorums should include the leader

Check whether $\log < \text{clientID}, \text{requestID} >$ is committed at position logSlotNum
 $\text{Committed}(\text{clientID}, \text{requestID}, \text{logSlotNum}) \triangleq$

Fast path

$$\vee \exists M \in \text{SUBSET } (\{m \in \text{messages} : \wedge \vee m.\text{mtype} = \text{MFastReply} \\ \vee m.\text{mtype} = \text{MSlowReply} \\ \wedge m.\text{logSlotNum} = \text{logSlotNum} \\ \wedge m.\text{dest} = \text{clientID} \\ \wedge m.\text{requestID} = \text{requestID}\}) :$$

Sent from a fast quorum

$$\wedge \{m.\text{sender} : m \in M\} \in \text{FastQuorums}$$

Matching view-id

$$\wedge \exists m1 \in M : \forall m2 \in M : m1.\text{viewID} = m2.\text{viewID}$$

One from the leader

$$\wedge \exists m \in M : m.\text{sender} = \text{Leader}(m.\text{viewID})$$

Hash values are consistent

$$\wedge \text{LET } \text{leaderReply} \triangleq \text{CHOOSE } m \in M : m.\text{sender} = \text{Leader}(m.\text{viewID})$$

IN

$$\forall m1 \in M : \text{IF } m1.\text{mtype} = \text{MFastReply} \text{ THEN } \\ m1.\text{hash} = \text{leaderReply}.\text{hash}$$

ELSE

TRUE SlowReply has consistent hash for sure

Slow path

$$\vee \exists M \in \text{SUBSET } (\{m \in \text{messages} : \wedge \vee m.\text{mtype} = \text{MSlowReply} \\ \vee \wedge m.\text{mtype} = \text{MFastReply} \quad \text{Leader only sends fast-reply} \\ \wedge m.\text{sender} = \text{Leader}(m.\text{viewID}) \\ \wedge m.\text{logSlotNum} = \text{logSlotNum} \\ \wedge m.\text{dest} = \text{clientID} \\ \wedge m.\text{requestID} = \text{requestID}\}) : \\ \wedge \{m.\text{sender} : m \in M\} \in \text{Quorums}$$

Matching view-id
 $\wedge \exists m1 \in M : \forall m2 \in M : m1.viewID = m2.viewID$
 One from the leader
 $\wedge \exists m \in M : m.sender = Leader(m.viewID)$

Check whether $log < clientID, requestID >$ is committed in view $viewID$
 $CommittedInView(clientID, requestID, viewID) \triangleq$

Fast path
 $\vee \exists M \in \text{SUBSET} (\{m \in messages : \wedge \vee m.mtype = MFastReply$
 $\vee m.mtype = MSlowReply$
 $\wedge m.dest = clientID$
 $\wedge m.requestID = requestID$
 $\wedge m.viewID = viewID\}) :$

Sent from a fast quorum
 $\wedge \{m.sender : m \in M\} \in FastQuorums$
 One from the leader
 $\wedge \exists m \in M : m.sender = Leader(m.viewID)$
 Hash values are the same
 $\wedge \text{LET}$
 $\quad leaderReply \triangleq \text{CHOOSE } m \in M : m.sender = Leader(m.viewID)$
 IN
 $\quad \forall m1 \in M : \text{IF } m1.mtype = MFastReply \text{ THEN}$
 $\quad \quad m1.hash = leaderReply.hash$
 $\quad \text{ELSE}$
 $\quad \quad \text{TRUE } \quad \text{SlowReply has consistent hash for sure}$

Slow path
 $\vee \exists M \in \text{SUBSET} (\{m \in messages : \wedge \vee m.mtype = MSlowReply$
 $\vee \wedge m.mtype = MFastReply \quad \text{Leader only sends fast-reply}$
 $\wedge m.sender = Leader(m.viewID)$
 $\wedge m.dest = clientID$
 $\wedge m.requestID = requestID$
 $\wedge m.viewID = viewID\}) :$

$\wedge \{m.sender : m \in M\} \in Quorums$
 Hash values are the same
 $\wedge \exists m1 \in M : \forall m2 \in M : m1.hash = m2.hash$
 One from the leader
 $\wedge \exists m \in M : m.sender = Leader(m.viewID)$

$SystemRecovered(viewID) \triangleq \wedge \exists RM \in \text{SUBSET} (Replicas) :$
 $\wedge Cardinality(RM) \geq QuorumSize$
 $\wedge \forall r \in RM : vLastNormView[r] \geq viewID$
 $\wedge \forall r \in RM : vReplicaStatus[r] = StNormal \quad \text{These replicas must be normal}$
 $\wedge vLastNormView[Leader(viewID)] \geq viewID$
 The leader of this view has also recovered or even goes beyond this view

Invariants

Durability: *Committed* Requests always survive failure

i.e. If a request is committed in one view, then it will remain committed in the higher views

One thing to note, the check of "committed" only happens when the system is still "normal"

While the system is under recovery (*i.e.* less than $f + 1$ replicas are normal),

the check of committed does not make sense

$Durability \triangleq \forall v1, v2 \in 1 \dots MaxViews :$

If a request is committed in lower view ($v1$),

it is impossible to make this request uncommitted in higher view ($v2$)

$\neg(\wedge v1 < v2$

To check *Durability* of request in higher views,

the system should have entered the higher views

$\wedge SystemRecovered(v2)$

$\wedge \exists c \in Clients :$

$\exists r \in 1 \dots MaxReqNum :$

$\wedge CommittedInView(c, r, v1)$

$\wedge \neg CommittedInView(c, r, v2))$

Consistency: *Committed* requests have the same history even after view changes

i.e. If a request is committed in a lower view ($v1$), then (based on *Durability* Property)

it remains committed in higher view ($v2$)

Consistency requires the history of the request (*i.e.* all the request before this request) remain the same

$Consistency \triangleq$

$\forall v1, v2 \in 1 \dots MaxViews :$

$\neg(\wedge v1 < v2$

To check *Consistency* of request in higher views,

the system should have entered the higher views

$\wedge SystemRecovered(v2)$

$\wedge \exists c \in Clients :$

$\exists r \in 1 \dots MaxReqNum :$

$\exists t \in 1 \dots MaxTime :$

Durability has been checked in another invariant

$\wedge CommittedInView(c, r, v1)$

$\wedge CommittedInView(c, r, v2)$

$\wedge LET$

$v1LeaderReply \triangleq \text{CHOOSE } m \in messages :$

$\wedge m.mtype = MFastReply$

$\wedge m.deadline = t$

$\wedge m.dest = c$

$\wedge m.requestID = r$

$\wedge m.viewID = v1$

$\wedge m.sender = Leader(v1)$

$v2LeaderReply \triangleq \text{CHOOSE } m \in messages :$

$\wedge m.mtype = MFastReply$

$\wedge m.deadline = t$

$$\begin{aligned}
& \wedge m.dest = c \\
& \wedge m.requestID = r \\
& \wedge m.viewID = v2 \\
& \wedge m.sender = Leader(v2)
\end{aligned}$$

$$\text{IN } v1LeaderReply.hash \neq v2LeaderReply.hash)$$

Linearizability: Only one request can be committed for a given position

i.e. If one request has committed at position i , then no contrary observation can be made

i.e. there cannot be a second request committed at the same position

Linearizability \triangleq

LET

$$\begin{aligned}
maxLogPosition & \triangleq Max(\{1\} \cup \\
& \{m.logSlotNum : m \in \{m \in messages : \\
& \quad \vee m.mtype = MFastReply \\
& \quad \vee m.mtype = MSlowReply\}\})
\end{aligned}$$

$$\begin{aligned}
\text{IN } & \neg(\exists c1, c2 \in Clients : \\
& \quad \exists r1, r2 \in 1 .. MaxReqNum : \\
& \quad \wedge \langle c1, r1 \rangle \neq \langle c2, r2 \rangle \\
& \quad \wedge \exists i \in (1 .. maxLogPosition) : \\
& \quad \quad \wedge Committed(c1, r1, i) \\
& \quad \quad \wedge Committed(c2, r2, i) \\
&)
\end{aligned}$$

Main Transition Function

Next \triangleq Handle Messages

$$\begin{aligned}
& \vee \exists m \in messages : \\
& \quad \wedge m.mtype = MClientRequest \\
& \quad \wedge m \notin vReplicaProcessed[m.dest] \\
& \quad \wedge HandleClientRequest(m.dest, m) \\
& \quad \wedge vReplicaProcessed' = \\
& \quad \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\
& \quad \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\
& \quad \wedge \text{UNCHANGED } vClientProcessed \\
& \quad \wedge DebugAction' = \langle \text{"HandleClientRequest"}, m \rangle
\end{aligned}$$

$$\begin{aligned}
& \vee \exists m \in messages : \\
& \quad \wedge m.mtype = MViewChangeReq \\
& \quad \wedge m \notin vReplicaProcessed[m.dest] \\
& \quad \wedge HandleViewChangeReq(m.dest, m) \\
& \quad \wedge vReplicaProcessed' = \\
& \quad \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\
& \quad \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\
& \quad \wedge \text{UNCHANGED } vClientProcessed
\end{aligned}$$

$$\wedge DebugAction' = \langle \text{"HandleViewChangeReq"}, m \rangle$$

$\forall \exists m \in messages :$

$$\begin{aligned} & \wedge m.mtype = MViewChange \\ & \wedge m \notin vReplicaProcessed[m.dest] \\ & \wedge HandleViewChange(m.dest, m) \\ & \wedge vReplicaProcessed' = \\ & \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\ & \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\ & \wedge \text{UNCHANGED } vClientProcessed \\ & \wedge DebugAction' = \langle \text{"HandleViewChange"}, m \rangle \end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned} & \wedge m.mtype = MStartView \\ & \wedge m \notin vReplicaProcessed[m.dest] \\ & \wedge HandleStartView(m.dest, m) \\ & \wedge vReplicaProcessed' = \\ & \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\ & \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\ & \wedge \text{UNCHANGED } vClientProcessed \\ & \wedge DebugAction' = \langle \text{"HandleStartView"}, m \rangle \end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned} & \wedge m.mtype = MSyncPrepare \\ & \wedge m \notin vReplicaProcessed[m.dest] \\ & \wedge HandleSyncPrepare(m.dest, m) \\ & \wedge vReplicaProcessed' = \\ & \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\ & \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\ & \wedge \text{UNCHANGED } vClientProcessed \\ & \wedge DebugAction' = \langle \text{"HandleSyncPrepare"}, m \rangle \end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned} & \wedge m.mtype = MSyncRep \\ & \wedge m \notin vReplicaProcessed[m.dest] \\ & \wedge HandleSyncRep(m.dest, m) \\ & \wedge vReplicaProcessed' = \\ & \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\ & \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\ & \wedge \text{UNCHANGED } vClientProcessed \\ & \wedge DebugAction' = \langle \text{"HandleSyncRep"}, m \rangle \end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned} & \wedge m.mtype = MSyncCommit \\ & \wedge m \notin vReplicaProcessed[m.dest] \\ & \wedge HandleSyncCommit(m.dest, m) \\ & \wedge vReplicaProcessed' = \end{aligned}$$

$$\begin{aligned}
& [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\
& \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\
& \wedge \text{UNCHANGED } vClientProcessed \\
& \wedge DebugAction' = \langle \text{"HandleSyncCommit"}, m \rangle
\end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned}
& \wedge m.mtype = MMissEntryRequest \\
& \wedge m \notin vReplicaProcessed[m.dest] \\
& \wedge HandleMissEntryRequest(m.dest, m) \\
& \wedge vReplicaProcessed' = \\
& \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\
& \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\
& \wedge \text{UNCHANGED } vClientProcessed \\
& \wedge DebugAction' = \langle \text{"HandleMissEntryRequest"}, m \rangle
\end{aligned}$$

$\forall \exists m \in messages :$

$$\begin{aligned}
& \wedge m.mtype = MMissEntryReply \\
& \wedge m \notin vReplicaProcessed[m.dest] \\
& \wedge HandleMissEntryReply(m.dest, m) \\
& \wedge vReplicaProcessed' = \\
& \quad [vReplicaProcessed \text{ EXCEPT } ![m.dest] = \\
& \quad \quad vReplicaProcessed[m.dest] \cup \{Msg2RLog(m, m.dest)\}] \\
& \wedge \text{UNCHANGED } vClientProcessed \\
& \wedge DebugAction' = \langle \text{"HandleMissEntryReply"}, m \rangle
\end{aligned}$$

Client Actions

$\forall \exists c \in Clients :$

$$\begin{aligned}
& \wedge vClientReqNum[c] < MaxReqNum \\
& \wedge ClientSendRequest(c) \\
& \wedge \text{UNCHANGED } \langle vReplicaProcessed, vClientProcessed \rangle \\
& \wedge DebugAction' = \langle \text{"ClientSendRequest"}, "" \rangle
\end{aligned}$$

Start Synchronization

$\forall \exists r \in Replicas :$

$$\begin{aligned}
& \wedge StartSync(r) \\
& \wedge \text{UNCHANGED } \langle vReplicaProcessed, vClientProcessed \rangle \\
& \wedge DebugAction' = \langle \text{"StartSync"}, "" \rangle
\end{aligned}$$

Replica Fail

$\forall \exists r \in Replicas :$

$$\begin{aligned}
& \wedge vReplicaStatus[r] = StNormal \\
& \wedge StartReplicaFail(r) \\
& \wedge \text{UNCHANGED } \langle vReplicaProcessed, vClientProcessed \rangle \\
& \wedge DebugAction' = \langle \text{"StartReplicaFail"}, "" \rangle
\end{aligned}$$

Leader Change

